

What needs to be done


Not a pile of models — a **fusion substrate** that joins every modality, with fast probabilistic producers hanging off it and a grounded query layer on top. The wet lab is the physical ground truth. Three views: the system, what it distills, and the order to build it.

HOW TO READ THE COLORS

- ☐ Data input
- ☐ Pipeline
- ☒ Fusion core · the spine
- ☐ Producer · probabilistic
- ☐ Query & interface
- ☐ Wet-lab ground truth
- ☐ Decision gate


The system

Four data modalities land in one core. Every producer **reads from and writes back to** the core as provenance-tagged results — so a score is never confused with a fact. The query layer and interface read the core; the wet lab validates the causal producer and supplies the organoid data.




What gets distilled

The payoff object. One variant accumulates every layer of evidence in a single traceable chain — today this is a seven-way join across seven people's separate outputs. Collapsing it into **one queryable record** is the feature that lets a researcher actually see.



The build, in order

Ownership decides everything, so it comes first. Then the spine. Then the **cheap high-leverage win** that runs on data you already have. Ambitious models come last, gated on wet-lab data — don't commit hiring to their timing until that data is real.



These three graphs mirror the written plan in `AfriCAN_DANCE_Build_Plan.md`. Phase 2 is the piece to ship first — it turns the engagement into a demonstrated result on existing data before any new sequencing.